

# A microarchitectural proposal for more aggressive exploitation of instruction level parallelism

D.A.D. Morano, D.R. Kaeli  
Northeastern University  
dmorano, kaeli@ece.neu.edu

18 July 2005

## Abstract

We present a new processor microarchitecture for managing aggressive parallel and speculative instruction execution. The goal is to explore ways to maximize processor performance for otherwise general purpose serial sequential program codes which do not lend themselves to explicit parallelization efforts. Instructions are fetched and dispatched for speculative execution as machine resources are available without first determining control or data dependencies. Rather, input dependencies are determined dynamically at execution time. Instructions remain in the processor (without being re-fetched or re-dispatched) and in a state of readiness for re-execution as correct input dependencies are determined. Committed instructions provide outputs to speculative instructions, which in turn re-execute as necessary in order to eventually converge on the correct committed program state. We present results showing some performance gains over more conventional processor microarchitectures with approximately equivalent hardware resources when executing integer sequential codes. Our proposed microarchitecture also features interconnection requirements that can be more naturally spatially separated than those of conventional superscalars. This lends the microarchitecture to a more distributed physical implementation, possibly allowing for larger physically scaled processors in the future.

## 1 Introduction

Although many high performance applications today can be parallelized at the application level and executed on tiled or clustered systems, there are and will continue to be requirements for achieving the highest performance on single threaded highly serially dependent program codes. We attempt to target this application problem space through the extraction of instruction level parallelism (ILP). The prospect of increased numbers of transistors available in silicon represents an opportunity to capitalize on the remaining ILP latent in even sequential (stubbornly non-parallelizable) programs. With limits to performance improvement through clock cycle reduction alone (witness also the recent flattening of higher clock frequencies), methods such as more aggressive ILP extraction in the microarchitecture become even more attractive.

Several studies into the limits of instruction level parallelism have shown that there is a significant amount of parallelism within typical sequentially oriented single-threaded programs (e.g., SpecInt-2000). The work of researchers including Lam and Wilson [9], Uht and Sindagi [18], Gonzalez and Gonzalez [4] have shown that there exists a great amount of instruction level parallelism that is not being exploited by any existing computer designs. Generally, ILP extraction is achieved by introducing multiple execution units into the microarchitecture and allowing each unit to operate as independently and as parallel as possible, yielding increased instructions per clock (IPC). One of the key problems with the addition of more parallel hardware resources is how to interconnect them in an efficient manner. Maintaining binary compatibility with existing instruction set architectures (ISA) is also generally a requirement, so we target this as well with our proposal.

Microarchitectures such as RAW [20, 15] or conventional cluster-based systems address the issue of parallelism within applications but only do so by exposing the spatially separated nature of their parallel-processor systems to the compiler and the application itself. This approach towards parallelism is an important one but can only address those applications that can be parallelized at a fairly coarse level.

Other microarchitectures that have employed the use of multiple execution units for ILP extraction are the Multiscalar-like processors [13, 14], the SuperThreaded processor model [17], and the Parallel Execution Window processor model [8]. The proposed MultiCluster machine model by Farkas et al. [3] are also in this category. Nagarajan also proposed a *Grid Architecture* of ALUs connected by an operand network [11]. However, the Multiscalar and Grid or grid-like microarchitectures also rely on the coordinated use of the compiler along with a new ISA. This differentiates these approaches with our own, which can be applied to any existing ISA. Although many proposed architectures and requirements for increased parallel programming of the application offer large potential performance increases over existing machines and non-parallelized programming styles, the requirement to provide high performance for existing (and future) non-parallelized codes with existing ISAs is likely a requirement for decades to come (withness the x86 ISA for example).

A microarchitecture that bears some similarity with our present proposal is the Ultrascalar design [5]. That proposal employed a relatively involved operand interconnection network (although logarithmic in scale to the size of the instruction window) that was switched according to instruction dependencies so as to route the appropriate input operand dependency to an instruction and that instruction’s output operand to the appropriate succeeding instructions. Our approach avoids the interconnection complexity of the Ultrascalar proposal and this is partly achieved through the dynamic determination of operand dependencies during execution using a common (not switched) operand forwarding interconnection fabric.

Our goal is a microarchitecture that features the benefits of speculative execution with relaxed control and data dependencies, such as done in the proposed Superspeculative microarchitecture [10], but with the ability to managing a large number of instructions simultaneously in flight. This makes the microarchitecture suitable for programs normally constrained by short dependency chains. An important distinction between our proposed microarchitecture and that of most others is that we dispatch instructions to special structures resembling reservation stations [1, 16], but instead of the instruction vacating the station upon instruction issue, it remains in the station until retirement. Instructions also dynamically determine their proper input dependencies, executing and re-executing as needed as new dependencies are determined. In our microarchitecture, the traditional reorder buffer (ROB) is eliminated along with the silicon layout routing complexity that is associated with it. Thus the layout density of implementing complex associative searches of the ROB for output result updates is avoided. These features are all achieved while also using moderately simple interconnections between our core machine components. In effect, the centralized functions of both operand dependency determination in the instruction window as well as the searching of the ROB for the most recent input operand is eliminated and instead distributed in our proposal. This allows for increased physical scalability that is not nearly as possible with the more conventional centralized instruction window and ROB-like approaches.

The rest of this paper is organized as follows. Section 2 gives an overview of the microarchitecture. Section 3 provides some further details on the more novel and core components of the microarchitecture. Section 4 provides some characterization and performance results for our microarchitecture through the simulated execution of benchmark programs. We summarize in Section 5.

## 2 Microarchitecture overview

### 2.1 Structural overview

The handling of main memory, and the cache hierarchy through the L1 instruction and L1 data caches are all conventional and similar to existing microarchitectures. Instruction fetch is also rather conventional. We also employ: a load-store-queue (LSQ) component, an architected register file (containing only committed registers), structures that closely resemble reservation stations or issue window slots, and rather conventional execution function units (FU). Our adaptation and enhancement of an issue window slot or reservation station is termed an *issue station* (IS). Figure 1 shows a high-level block diagram of our microarchitecture showing the major instruction execution components. The memory hierarchy, as well as details of the instruction fetch unit, are not shown as they are similar to those of existing machines. Although there is only a single load-store-queue (LSQ) and an architected register file, the numbers of all other components can vary with a specific machine implementation. All ISs are identical, allowing for any instruction to be dispatched to each. Function units can be duplicated and functionally differentiated (meaning different types of units would handle different classes of instructions) as desired.




Figure 1: *High-level block diagram of our microarchitecture.* Issues stations (IS) are shown on the left and various function units (FU) on the right. An architected register file and a load-store-queue are shown at the top. Bidirectional operand request and forwarding buses are shown vertically oriented (to the right of the Issue Stations). Buses to transport an instruction operation and its source operands to the function units are also shown. Likewise buses to return result operands are present.

Like other machines with reservation stations, we dispatch decoded instructions from the fetch unit to these ISs when one or more of them are empty (available for dispatch) or becoming empty on the next clock cycle. Instructions are dispatched in-order. The number of instructions dispatched in any given clock cycle is the lesser of the number of ISs available and the dispatch width (as counted in numbers of instructions). Instructions are always dispatched to ISs when one or more are free (or becoming free) and are generally dispatched without first knowing their input dependencies or having those corresponding operands available.

The vertical buses, roughly in the center of Figure 1, labeled *operand forward-store* and *operand backward-request*, are bidirectional and multi-master bus groups that allow for operands to be requested from and forwarded to ISs. A bus group is a set of one or more buses in parallel (the exact number of which is an implementation option) that can be used to increase transfer bandwidth through alternative paths. Bus paths that do not start and end at the same repeater junctions are also possible and have been explored. These operand buses (forward-store and backward-request) provide the means by which instructions acquire their necessary operands for proper execution. Requests for operands are generated by ISs and are placed on the backward-request bus group, while answers to operand requests are placed on the forward-store bus group. Operand requests also travel back to the register file and LSQ (each also attached to these buses), and those units can likewise forward operands in response to requests. Operand switching interconnects other than a bus arrangement are also possible. Operands only have to traverse the operand interconnect arrangement (whatever it might be) in a single direction ; namely, toward younger dispatched instructions for forwarded operands and toward older dispatched instructions for backwarded <sup>1</sup>. This allows additional flexibility in both the choice of an interconnection fabric and the ability to electrically isolate segments of it due to silicon layout propagation delay constraints.

However, the simplicity of parallel buses as the basic operand interconnect generally makes for simpler silicon hardware layout and better future physical scalability when used in conjunction with some sort of repeater mechanism. Our microarchitecture allows for long buses (longer than the given silicon constraints and propagation delays) to be split using repeater like units since it does not assume any fixed delays (in clocks) for the transfer of operands. This also effects a great deal of flexibility for both bus and non-bus interconnection arrangements. The details of how operand bus repeaters work is not discussed further in this present paper.

Buses are also provided for bringing instruction codes and operands from the issue stations to the FUs and back again. These are also arranged in a parallel group, the basic group width being determined by the desired issue width to the FUs. This arrangement is not too dissimilar to that between the issue window and the FUs of some conventional microarchitectures.

Collectively, all of the components discussed in this section (and shown in Figure 1) are termed the

<sup>1</sup>Perhaps awkward, but consistent with its opposite of being forwarded.

*execution window.*

## 2.2 Operational overview

Instructions are decoded after fetch and stored in a buffer (in decoded form) for possible dispatch to the ISs. As instructions are dispatch to the ISs, a *time-tag* is assigned along with the decoded instruction. The time-tag is a small positive integer, large enough to represent the number of ISs in a particular implementation. New time-tags are assigned sequentially higher values starting from zero and all time-tags in the execution window are decremented as instructions commit (effectively decremented by one for each committed instruction). This is how time-tag value wrap-around is handled. This arrangement always keeps higher valued time-tags representing younger dispatched instructions. This property is used during the operand snoop process (described in detail later). Note that the whole set of ISs effectively serve the function of the ROB in more conventional machines, albeit in a more distributed way.

Once instructions are dispatched to ISs, they issue requests on the backward-request bus group for their input operands. The ISs then wait for plausible (generally speculative) input operands through the process of snooping the operand forward-store bus group. The process of snooping for operands consists of checking the output operands created by other ISs (or forwarded from the architected register file) to see if they are feasible candidates to satisfy input dependencies. Feasible input operands are any of those that have the same architected address. For register operands, this is the architected number of the register. For memory operands it is the architected memory address. These feasible operands may be from either the closest older instruction that created a feasible output operand or it may be from still older instructions. In the former case the operand may be speculative or not (depending on whether the instruction ready to commit or not), but in the latter case the operand will generally be speculative but could be the same value as the proper non-speculative operand by coincidence. The details of the operand snooping process is covered in greater detail in a subsequent section. Issue stations also snoop the backward-request bus group to see if they have an output operand that may satisfy an outstanding operand request from a younger IS (determined by time-tag value). Note that both the LSQ and the architected register file also snoop for operand requests (although each only for its particular operand type). In this way all operand dependencies are determined dynamically during and intermixed with execution or re-execution. Note that although operands received from the LSQ or architected register file are always valid (having been committed), operands received from older ISs are generally still speculative.

Once feasible input dependencies have been acquired, the IS can execute its currently associated instruction. Some instructions are executed within the IS station itself while the remaining ones need to contend with other ISs for FU availability. In our present work, the instructions that execute within the ISs are those with an instruction type of: load-store or control-flow-change. With this decision, the logic and state complexity of both the ISs and the FUs is substantially reduced over the case where all instructions must proceed through FUs. Note that typically there is a full add function required by many or most load-store instructions in order to perform an address calculation (and the equivalent for comparisons in branch instructions). Although this *add* function could be carried out in a FU, we have chosen to place the adder inside the IS. Although this represents a sizeable silicon resource enhancement (substantially more transistors) to the IS, it seems a reasonable decision based on the increasing availability of transistors with newer technology. However, this design tradeoff would likely not hold for any CISC-like ISAs where these instructions are more complex than their typical RISC definitions. In general, the present arrangement represents a division of execution resources between both the ISs and the FUs, effectively increasing issue width. The process of winning a FU execution-slot constitutes an instruction issue, in the more conventional sense. All instruction issues and executions in the machine are entirely out of order.

In the case of requiring an FU for execution, the output results are transferred back to the IS before being forwarded further. In both execution type cases, results of instruction executions are retained within the IS until retirement. Forwarding of instruction execution results to younger dispatched instructions is always done from the IS itself, rather than from the output of the FUs (as is often the case in most machines). This is done since the ISs are the place where the results are stored as tentative values. This is consistent with the idea that the ISs are serving the function of the ROB, which would otherwise store tentative results. This also eliminates the requirement for the ISs to snoop the outputs of the FUs for possible new input

dependency candidates (dramatically reducing interconnection requirements).

What has been implicit in the discussion so far is that whenever operands are snooped by the ISs and are found to be feasible input dependencies, a re-execution of the associated instruction may be triggered. Re-executions do not have to occur if the values of the input dependencies do not change. This re-execution technique is somewhat analogous to the reply mechanisms being currently proposed, where an instruction is retained in the issue window and re-issued if necessary. However, our proposal presents this idea as more fundamental to the whole microarchitecture and more natural as the required input operands are immediately present with (located adjacent to) the instruction to be re-executed (replayed).

Instruction commitment can occur on each clock, is in-order, and proceeds from the oldest program-ordered instruction through younger instructions until an instruction is reached that doesn't meet the requirements for commitment. Preference to win a FU execution-slot by an issue station is given to the ISs holding the oldest dispatched instructions (this facilitates movement towards commitment). Instructions can only commit if they have executed at least once, have not received a new feasible input operand (which could trigger re-execution and a different committed result), and are finished forwarding output operands to younger dispatched instructions. An output operand is considered forwarded when it is placed on the forward-store bus group. The last constraint guarantees that all younger instructions can process the final output operands of older instructions before those older instructions commit and are therefore removed from the execution window. Committed outputs are also forwarded to the architected register file or LSQ (as appropriate for operand type, register or memory). Note that instruction commitment is not required to be delayed so as to wait for verification that younger instructions have received their new inputs from the committing instruction. The only requirement for the committing instruction is that any of their output operands not already forwarded win a transfer slot on the forward-store bus group. This guarantees that the operand will be seen by possibly dependent instructions.

When a branch is resolved (ready for commitment), instructions in ISs beyond the resolved branch are abandoned if they are not on the resolved branch path. This is similar to what would occur in a microarchitecture with both an issue window and an ROB, except that both actions occur in our ISs as opposed to the corresponding action occurring in both the issue window and the ROB. Program interrupts and exceptions are handled similarly as they would in a machine with an ROB.

## 3 Core component detail

### 3.1 Issue Stations

The ISs provide the most significant distinction of this microarchitecture from most others. These are similar to reservation stations but contain additional state and logic that allows for dynamic operand dependency determination as well as for holding a dispatched instruction (its decoded form) until it is ready to be retired. There is state inside the station that is relevant to the instruction itself and specific to the operands of that instruction (both source and destination operands).

The state that is primarily associated with the instruction itself consists of:

- instruction address
- instruction operation
- execution state
- time-tag
- instruction predication information

The *instruction operation* is derived from the decoded instruction and specifies the instruction class and other details needed for the execution of the instruction. This information may consist of subfields and is generally ISA specific. The *instruction address* and *predicate* state are only used when dynamic predication [19] is done within the microarchitecture. The time-tag value is used to order this instruction with respect to all others that are currently within the execution window of the machine. The time-tag is also used as part of the operand snooping logic (discussed more later). The *execution state* value constitutes the state used by various state machines within the IS for controlling its operation and for determining readiness for commitment.

The remainder of the state consists of one or more input source operands and one or more output destination operands. All operands regardless of type and whether source or destination occupy a similar structure

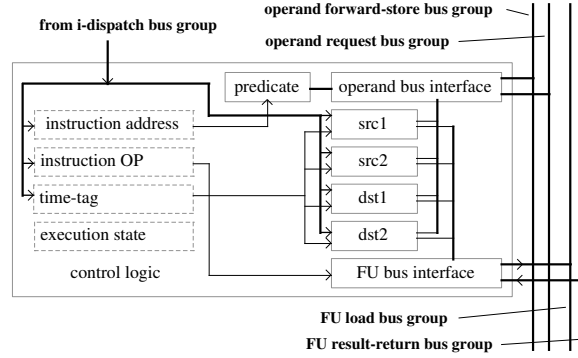


Figure 2: *High-level block diagram of our Issue Station.* The major state and sub-blocks associated with an Issue Station is shown. General instruction state (shown in dash-lined boxes) is on the left, while four operand blocks (two source and two destination), and the four primary bus group interfaces (grouped by function at upper right and lower right) to the rest of the execution window are on the right.

within an IS, termed an *operand block*. More detail on these operand blocks and operand management is provided in the next section.

A simplified block diagram of our IS is shown in Figure 2. In this example, a total of four operand blocks are shown, labeled: *src1*, *src2*, *dst1*, and *dst2*. The number of source and destination operand blocks that are used for any given machine is dependent upon the requirements of the ISA implemented.

### 3.2 Operands

The types of operands are distinguished: register and memory. Operands blocks are constructed to hold either type. The state within an operand block consists of:

- type of operand
- time-tag
- address
- size
- previous value
- value

The operand *time-tag* serves an analogous purpose as the time-tag register within an IS, except that it applies specifically to this particular operand rather than to the instruction as a whole. Again, this time-tag is used in the operand snooping logic and allowed for the dynamic discovery of dependencies for instructions.

The *address* field differs depending on the type of the operand. For register operands, the address would be the name of the architected register. All ISA architected registers are typically provided a unique numerical address. These would include the general purpose registers, any status or other non-general purpose registers, and any possible ISA (architected) predicate registers (like those in the iA-64 ISA [7, 12]. For memory operands, the identifying address is just the programmer-visible architected memory address of the corresponding memory value.

The *size* is only used for memory operands and holds its size in bytes. The *value* holds the present value of the operand calculated from this present instruction (if it has executed at least once). The *previous value* is only used for destination operands and holds the value that the operand had before it may have been changed by the execution of the present instruction. The previous value is used when a forwarded operand with a specific address was incorrect. This situation occurs when addresses for memory operands are speculatively calculated but are later determined to have changed. An operand with the old address is forwarded with the previous value to correct the situation.

Figure 3 shows a simplified block diagram of an operand block along with its major data paths for its major functions. These major functions consist of:

- snooping for feasible input operands
- snooping for operand requests from elsewhere

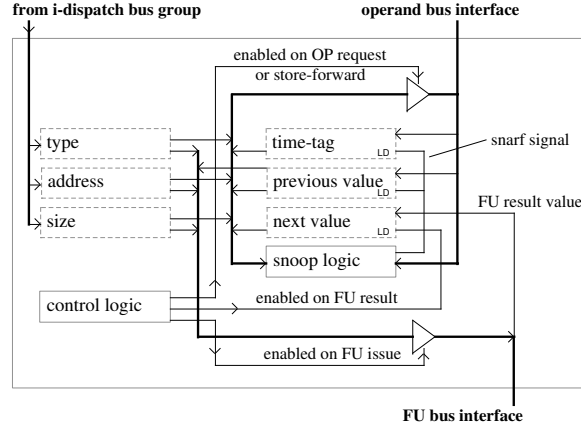


Figure 3: *Block diagram of an Operand Block.* Each Operand Block holds an effectively renamed operand within the Issue Stations. Several operand blocks are employed within each Issue Station depending on the needs of the ISA being implemented. The primary state register information maintained for each operand (shown in dash-lined boxes) along with the major data paths and enabling signals for its major functions are shown.

- issuing instruction code to FUs (if required due to instruction type)
- receiving FU store-results back (if required due to instruction type)

In effect, a full renaming of all operands is realized for all instructions in flight in the machine. All false dependencies are thusly avoided. Full names for operands consist of the components:

- type of operand
- time-tag
- address

and these components, taken together, fully disambiguate all in-flight operands from each other (implementing full renaming).

### 3.3 Operand forwarding and snooping

Operands resulting from the execution of instructions are transmitted forward, using the forward-store bus group, or use by younger (in program order) waiting instructions. All forwarded operands are snooped by the operand blocks within those ISs containing younger dispatched instructions. The information associated with each operand that is forwarded is referred to as a *transaction* and consists of:

- transaction type
- operand type
- address
- time-tag of the originating issue station
- data value for this operand

The time-tag forwarded with the operand is that of the originating IS (instruction instance). The operand information above is typical of both register and memory operand transactions and the use of the *operand type* distinguishes one from the other. The *transaction type* field is used to designate whether the transaction represents a store from a previous instruction or an indication that a previously forwarded operand is no longer valid.

When a set of matching conditions is found by examining the component parts of the operand while snooping, an acquisition of the operand is then effected. That acquisition is termed a *snarf*. A snarf for a particular operand within an IS occurs when: the operand type and address of the snooped operand match that of the stored operand, and the snooped time-tag is both less than the current instruction time-tag (stored in the IS) and is less than or equal to the last time-tag snarfed for the given stored operand. In the case of a snarf, the stored operand time-tag (TT) and previous-value (PV) registers are reloaded with the associated fields from the snooped operand transaction. Additionally, if the snooped operand data value is




Figure 4: *Snooping logic for operand updates.* The snooping logic for one of several possible source operands is shown. This logic would reside in each of the operand blocks within an IS and they would all perform the snoop operation simultaneously. Just one operand forwarding bus is shown being snooped but typically several forwarding buses are snooped simultaneously.

different than the stored operand previous-value, an execution or a re-execution is scheduled for the current IS. However, if the snooped data value is the same as the stored previous-value, no new execution is triggered. This eliminates some unnecessary re-executions.

A simplified schematic diagram of the logic used for operand snooping is shown in Figure 4. The *time-tag* and *previous-value* registers within an operand block are reloaded with new values on each snarf, while the *instruction time-tag* register in the IS is only loaded when an instruction is dispatched. The operand block *address* register is either loaded at instruction dispatch or may be loaded during instruction execution for some instructions (for example by load-store instructions). Note also that very similar snooping logic is located in both the LSQ and the architected register file allowing them to discover operand requests so that they can likewise respond as ISs do.

## 4 Experimental results

In this section we present a first look at some experimental results from simulation of the machine presented. We designed a simulator to evaluate the performance and characteristics of our proposed design. The simulator models the major machine components mostly behaviorally but with the same state and state transitions that would be in actual hardware. Components are then interconnected structurally. The simulator implements the Alpha ISA.

We use ten of Spec2000 integer benchmarks (listed in Table 2 below) to evaluate the potential of this microarchitecture<sup>2</sup>. For all simulations, the initialization phases of all programs are skipped using a fast-forward mechanism. This allows the subsequent functional simulation (which takes the real bulk of simulation time) to operate on the most characteristic part of the benchmarks. After the initial fast-forward operation, a phase of one million instructions are executed to warm up machine components that have longer state residency times. This currently includes the cache hierarchy (L1 and L2 caches) and the branch predictor. Then a short sequence of instructions are executed to prime machine components such as the ISs. This sequence is approximately equal to two times the number of ISs configured for the target machine. Finally, we execute on the main functional cycle simulator for the next 100 million instructions.

For all simulations, we used separate instruction and data L1 caches, but a unified L2 cache. The data caches use a write-back policy with a least recently used block replacement algorithm. Other configuration parameters of the machine for the simulations are shown in Table 1.

We present the IPC results of our proposed microarchitecture with that of a baseline superscalar microarchitecture that is similarly configured. We used the SimpleScalar MASE framework [2] to simulate a conventional superscalar (approximately a MIPS R10000 in the case of SimpleScalar MASE).

The baseline superscalar machine includes an instruction window consisting of reservation stations and a re-order buffer (ROB) to store speculative execution result registers pending commitment. Instructions for

<sup>2</sup>We have not yet adequately handled the system calls in the remaining two programs of the SpecInt-2000 suite.



Table 1: *General machine characteristics.* These machine parameters are used for all simulations unless otherwise specified.

L1 I cache access latency	1 clock
L1 I cache size	32 KBytes
L1 I block size	32 bytes
L1 I organization	direct mapped
L1 D cache access latency	2 clocks
L1 D cache size	128 KBytes
L1 D block size	32 bytes
L1 D organization	2-way set assoc.
L2 cache access latency	20 clocks
L2 cache size	1 MBytes
L2 block size	32 bytes
L2 organization	4-way set assoc.
main memory access latency	150 clocks
branch predictor	2-level w/ XOR
	16k PHT entries
	8 history bits
	32k BHT entries
	sat. 2-bit counter
fetch width	8 instructions
FU issue width	4 operations
issue width	4
number integer FUs	4
number other FUs	1 each
forwarding buses	8
bus traversal latency	1 clock
integer FU latency	1 clock
other FU latencies	3 to 17 clocks

the baseline superscalar are fetched and dispatched to the instruction window where they wait for input data dependencies to become ready. When all instruction input dependencies are ready, and as issue bandwidth allows, instructions are issued to the function-unit pipelines. Register results are stored in the ROB until commitment. Both the baseline superscalar and our proposed microarchitecture flush the execution window on a resolved mispredicted conditional branch. This is a fairly typical superscalar execution arrangement and is fixed in the Simplescalar MASE simulator.

Although an exact comparison of the two machines is not possible due to their very different construction, we have arranged for both the baseline machine and our proposal to have either an exact or a very close correspondence in the amount and number of hardware resources. Both machines are configured with identical cache arrangements and cache configurations. Both also employ the same branch predictor and predictor configuration. Both also implement a four-wide issue machine. For both machines, the issue width is the maximum number of instructions that can be issued in a single clock cycle. However the baseline superscalar employs reservation stations and an ROB while our microarchitecture uses our novel ISs. We therefore roughly equate the number of ISs of the proposed machine with the combination of both the number of instruction window slots (reservation stations) and ROB entries of the baseline superscalar. Our results are for 128 ISs in the proposed microarchitecture and both 128 issue window slots (reservation stations) and 128 ROB entries for the baseline superscalar. This represents a modest sized machine of today.

The IPC results for the baseline superscalar and our proposed machine are shown in columns two and three of Table 2 respectively. In order to get a feel for how much re-execution of instructions occurs, we also list in the results (in column 4 labeled REX) the percent additional instructions executed as compared with the number of committed instructions. We also calculated the harmonic mean of the IPC across all benchmarks (last entry in table). Columns 5 is discussed later. Comparing the harmonic mean IPC results in columns two and three of the Table 2 (last entry), our proposed machine attained a speedup (1.61 divided by 1.10) of approximately 1.46 as compared with the baseline superscalar. This was achieved using approximately the same amount of hardware resources. All individual programs performed better with the exception of the BIPZ2 program. We are still exploring why BZIP2 performed poorly on our machine as compared with the baseline superscalar. We have not found any single internal metric that seems to present a particular bottleneck for BZIP2, but it also may be that some feature of the conventional machine allowed it to perform unusually well for some reason. Column 4 of Table 2 (titled REX) shows the percentage

Table 2: *IPC and re-execution results.* IPC performance of a MASE baseline superscalar machine and our proposed microarchitecture is presented. The column titled REX gives the percent of extra instructions executed for the newly proposed machine as compared with the committed instructions.

	baseline	new		new-extra
	IPC	IPC	REX	IPC
bzip2	2.11	1.98	97.8%	2.41
crafty	1.34	1.96	79.2%	2.33
eon	1.25	2.62	110.4%	3.12
gcc	1.32	1.70	96.2%	2.41
gzip	1.35	1.42	98.0%	1.93
parser	0.80	1.23	115.8%	1.41
perlbnk	0.64	1.44	92.3%	1.52
twolf	1.16	1.32	88.2%	1.60
vortex	1.05	2.61	103.7%	3.75
vpr	1.08	1.13	96.1%	1.39
H-MEAN	1.10	1.61		1.97

of committed instructions that incurred re-executions. Most benchmarks exhibit a behavior of executing somewhere around 200% (within about plus or minus 20%) of the committed number of instructions. This compares similarly to the amount of execution needed in proposals like the SlipStream processor [6], except that neither two threads nor two processors need be dedicated to the execution of a single program, as it done in that processor. Rather, approximately two times the execution was performed within the resources of a single core and threaded machine.

Column 5 provides the IPC results for a version of the machine where additional instructions are executed within the ISs as compared with the machine of column 3. Minimally, control-flow and memory load-store instructions are executed in the ISs. However with a very small additional amount of silicon, simple bit-logic instructions and integer add and subtract instructions (including integer compare) can also be executed within the ISs. This arrangement (transistors permitting) increases the amount of resources available for executions, thus increasing parallelism and IPC performance. This yields approximately 22% better IPC performance than with the minimal configuration proposed machine, and approximately 79% IPC improvement of the baseline machine.

We note that with a higher IPC, power consumption can be reduced while achieving the same performance as the baseline machine by lowering clock frequency. Future research might also realize power savings through the possible elimination of redundant re-executions.

## 5 Summary

We have described a new microarchitecture that allows for both control and data speculative execution, but also does so in a way where necessary re-executions are handled quickly and cheaply in the hardware, without requiring either re-fetch or re-dispatch. The necessary and complicated instruction dependency enforcement is achieved dynamically during execution (and re-execution) using time-ordering tags that maintain relative program order of instructions and all operands in flight. Binary program compatibility with existing ISAs (an important feature in the market place) is also maintained with our proposal. Our results show that our proposed machine achieves approximately 46% better IPC performance over a conventional machine of roughly equivalent silicon resources, and approximately a 79% IPC improvement given some modest additional silicon to facilitate executing additional instructions in the issue stations. Physical scalability is possibly facilitated due to the distributed handling of operand dependency determination and forwarding.

## References

- [1] Anderson D., Sparacio F. and Tomasulo F. The IBM/360 Model 91: Machine Philosophy and Instruction Handling. *IBM Journal*, 11(1):8–24, Jan. 1967.
- [2] Austin T.M. and Burger D. SimpleScalar Tutorial. In *Proceedings of the 30th International Symposium on Microarchitecture*, Nov 1997.

- [3] Farkas K.I., Chow P., Jouppi N.P., Vranesic Z. The multicluster architecture: Reducing cycle time through partitioning. In *Proceedings of the 30th International Symposium on Microarchitecture*, pages 149–159, 1997.
- [4] Gonzalez J. and Gonzalez A. Limits on Instruction-Level Parallelism with Data Speculation. Technical Report UPC-DAC-1997-34, UPC, Barcelona Spain, 1997.
- [5] Henry D.S., Kuszmaul B.C., Viswanath V. The ultrascalar processor – an asymptotically scalable superscalar microarchitecture. In *Proceedings of the 20th Anniversary Conference on Advanced Research in VLSI*. IEEE, Mar 1999.
- [6] Ibrahim K.Z., Byrd G.T., Rotenberg E. Slipstream execution mode for CMP-Based multiprocessors. In *Proceedings of the 9th International Symposium on High Performance Computer Architecture*. IEEE, Feb 2003.
- [7] Intel Corp. *iA-64 Application Developer's Architecture Guide*, 1999.
- [8] Kemp G.A., Franklin M. PEWs: A decentralized dynamic scheduler for ILP processing. In *Proceedings of the 24th International Conference on Parallel Computing*, pages 239–246, 1996.
- [9] Lam M.S. and Wilson R.P. Limits of Control Flow on Parallelism. In *Proceedings of the 19th International Symposium on Computer Architecture (ISCA)*, pages 46–57. ACM, May 1992.
- [10] Lipasti M.H. and Shen J.P. Superspeculative Microarchitecture for Beyond AD 2000. *IEEE Computer*, 30(9), Sep 1997.
- [11] Nagarajan R., Sankaralingam K., Burger D. and Keckler S.W. A design space evaluation of grid processor architectures. In *Proceedings of the 34th International Symposium on Microarchitecture*, New York, NY, Nov 2001. ACM Press.
- [12] Schlansker M.S. and Rau B.R. EPIC: Explicitly parallel instruction computing. *Computer*, 33(2):37–45, Feb 2000.
- [13] Sohi G.S., Breach S. and Vijaykumar T.N. Multiscalar Processors. In *Proceedings of the 22th International Symposium on Computer Architecture*, New York, NY, Jun 1995. ACM Press.
- [14] Sundararaman K.K., Franklin M. Multiscalar execution along a single flow of control. In *Proceedings of the International Conference on Parallel Computing*, pages 106–113, 1997.
- [15] Taylor M.B., Kim J., Miller J., Wentzlaff D., Ghodrat F., Greenwald B., Hoffmann H., Johnson J., Lee J., Lee W., Ma A., Saraf A., Seneski M., Shnidman N., Strumpen V., Frank M., Amarasinghe S. and Agarwal A. The RAW microprocessor: A computational fabric for software circuits and general purpose programs. *IEEE Micro*, 2002.
- [16] Tomasulo R.M. An Efficient Algorithm for Exploiting Multiple Arithmetic Units. *IBM Journal of Research and Development*, 11(1):25–33, Jan 1967.
- [17] Tsai J-Y., Yew P-C. The superthreaded architecture: Thread pipelining with run-time data dependence checking and control speculation. In *Proceedings of the International Conference on Parallel Architectures and Compilation Techniques*, pages 35–46, 1996.
- [18] Uht A. K. and Sindagi V. Disjoint Eager Execution: An Optimal Form of Speculative Execution. In *Proc. MICRO-28*, pages 313–325. ACM, Nov 1995.
- [19] undisclosed2. undisclosed2. In *undisclosed*, Mar 2002.
- [20] Waingold E., Taylor M., Srikrishna D., Sarkar V., Lee W., Lee V., Kim J., Frank M., Finch P., Barua R., Babb J., Amarasinghe S., and Agarwal A. Baring it all to software: RAW machines. *IEEE Computer*, 30(9):86–93, 1997.