

Supplemental Data for Characterization of Register Temporal Locality

David A. Morano
Northeastern University
dmorano@ece.neu.edu

25 October 2002

Abstract

This paper is a supplement to a paper submitted to ISPASS'03. In that paper various intervals associated with register and memory accesses were explored. This paper primarily provides additional data that was not presented in that paper. Detail about the three register access intervals for each of the benchmark programs examined is presented. Additionally, new data is presented for all of the benchmarks on the read and write access frequencies for the registers by register address.

1 Introduction

This paper provides supplementary data on the access intervals for registers and memory locations over what has been previously reported in the paper *Implications of Register and Memory Temporal Locality for Distributed Microarchitectures* by Khalafi, et al [1]. In that paper, data for register access intervals was provided in an abbreviated form. In this paper, we present the data again but separated out for each of the ten benchmark programs that were investigated. Further, we present additional data about register read and write access frequencies per register address that was not presented. We only present new data for register access intervals. A large set of data for the memory variable access intervals was already provided in the preceding paper.

The rest of this paper is organized as follows. In Section 2 provides an additional example to illustrate how the various access intervals are calculated. Section 3 presents our additional characterization results. In Section 3 we discuss some of the data results. Section 4 summarizes the present work.

2 Example Interval Calculation

A simple and contrived code example illustrating the meaning and method of calculation of the three intervals, that we have defined, is shown in Figure 1. Unlike the previous example, this one features a loop. In this

Table 1: *Simple code example illustrating the different types of access intervals.* Alongside the instructions are the particular events generated from those instructions and any intervals that can be calculated as a result.

label	instruction	event	intervals determined
I1	r1 <= c1	def(r1)	
I2	r2 <= r1 + c2	def(r2), use(r1)	access-use(r1)=1, def-use(r1)=varies
I3	beq r2, I2	use(r2)	access-use(r2)=1, def-use(r2)=1
I4	r1 <= c3	def(r2)	useful-lifetime(r1)=2001

Table 2: *Access Interval Statics on the Simple Code Example of Table 1.* The mean and standard deviation is given for access-use, useful-lifetime, and def-use intervals.

interval	mean	standard deviation
access-use	1.5	0.5
useful-lifetime	3.0	63.2
def-use	500.5	645.1

simple code example, the terms *c1* and *c2* are arbitrary immediate constants encoded within the instructions. We assume that this code loops 1000 times and then control falls through the conditional branch at *I3*. More precisely, instructions *I2* and *I3* execute 1000 times each. This snippet is interesting because it shows how the various access intervals can sometimes vary widely in their average behavior from each other when just simple code constructs are encountered.

For all three access intervals we are investigating, their statistics are given in Table 2 for this simple code example. In particular, note the widely varying means for the useful-lifetime (3.0) and the def-use (500.5) intervals. As might be expected, the access-use interval is shorter (on average) than the other two but it may be somewhat interesting to see the difference between the other two interval averages. In each case, the standard deviation gives a clue to the amount of dispersal in the data. Instruction *I2* is the cause of the def-use interval dispersal. While the useful-lifetime of all registers (when averaged together) is kept low due to the presence of the def in instruction *I2* even the def in instruction *I4* does not bring the useful-life average over about 3.0. In fact, without the def of instruction *I4*, the average useful-lifetime of all registers would have stayed at about 1.0 with a standard deviation of about 0.0. The key feature of this code example is the use (by the compiler or a human programmer) of setting up the invariant variable in register *r1* prior to entering the loop. Hopefully, this would occur commonly in optimizing compilers. Of course, real code contains more interesting cases than shown in this example but this illustrates how the various intervals can vary from each other due to various code factors.

Not illustrated here but which occurs not uncommonly is that writes to variables (whether registers or memory locations) can be abandoned due to a change in control flow from a conditional branch. In these sorts of cases, the compiler may have setup some variable expecting the control flow to proceed in a certain direction, which is subsequently thwarted by the actual execution.

3 Characterization Results

Here we present data for both the three access intervals presented previously but with the full set of data shown for each benchmark program. In addition to the access intervals, we also show register read and write frequencies. This data has not been previously presented in any form.

3.1 Register Access Interval Results

In this section, we show the register access interval results for our ten benchmark programs. Unlike the previous paper, we show the data collected for each program for individual benchmark programs rather than overlaid on a single graph. Data for register access-use, useful-lifetime, and def-use intervals are shown in turn. The data for register access-use intervals are shown in Figures 1 and 2. Results from benchmark programs BZIP2, COMPRESS, CRAFTY, GCC, and GO are shown in Figure 1 while the results from programs GZIP, IJPEG, MCF, PARSER, and VORTEX are shown in Figure 2.

The data for register def-last-use (or useful-lifetime) intervals are shown in Figures 3 and 4.

The data for register def-use intervals are shown in Figures 5 and 6.

Figure 7 shows (top to bottom) the three types of access intervals: access-use, useful-lifetime, and def-use. For each type of access interval, the data for each of the ten benchmark programs are overlaid on the same graph. For each access interval type, both a density and a distribution is provided. This presentation of the register interval access data has been already presented but is repeated here for the reader’s convenience. As

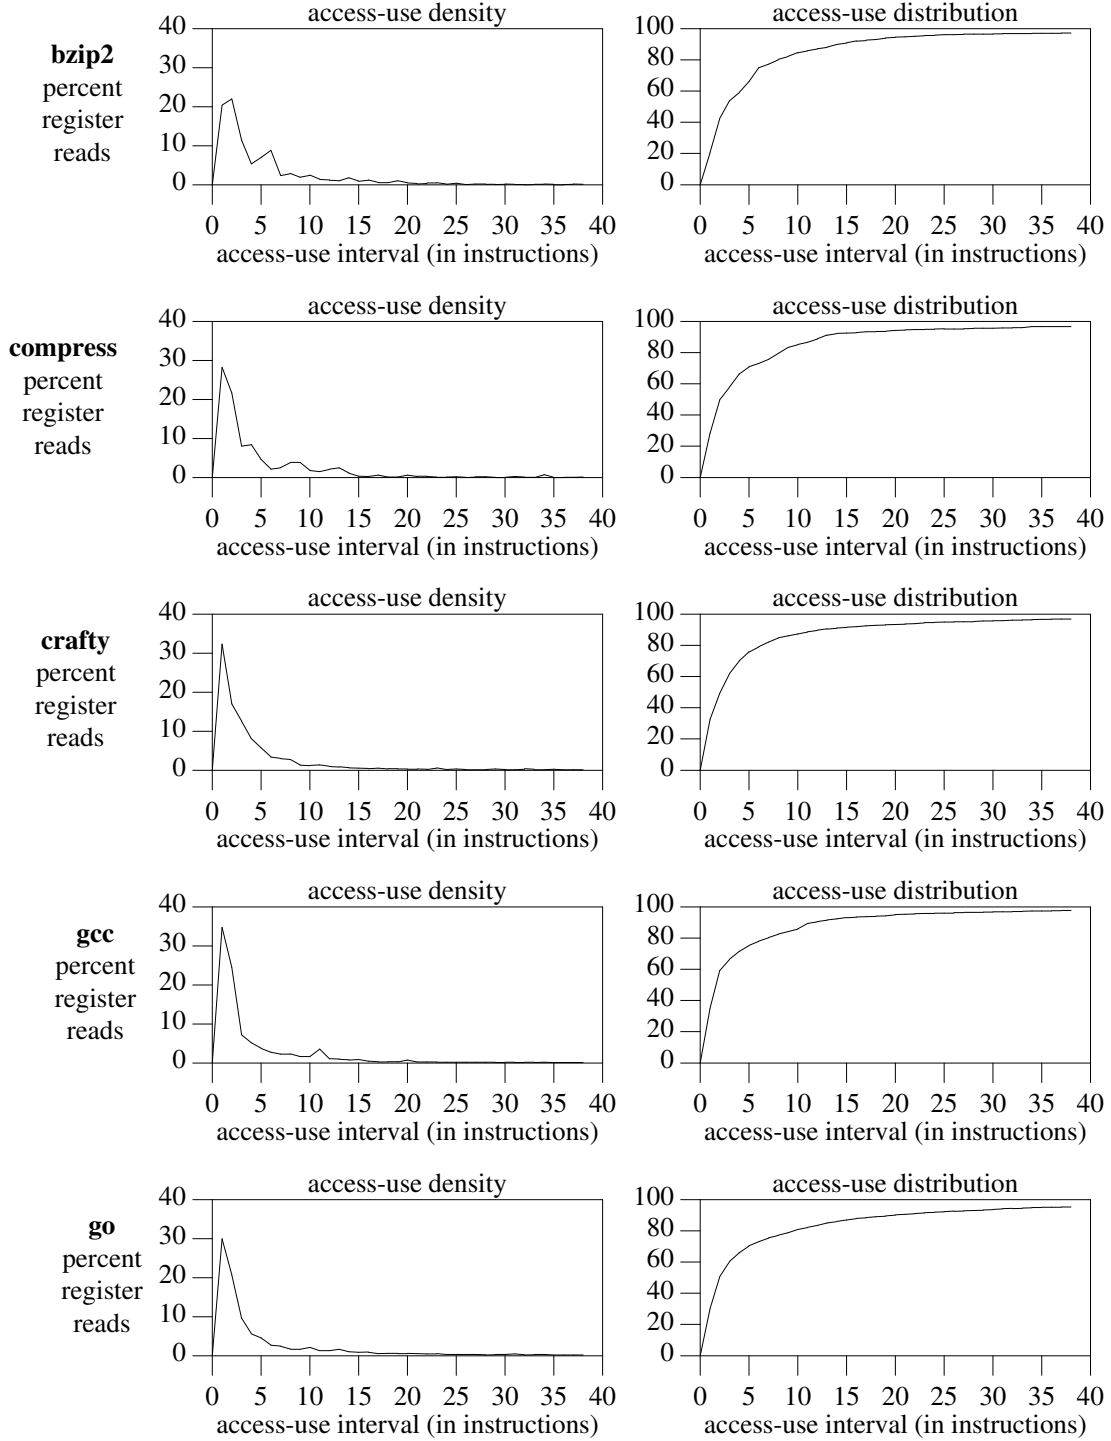


Figure 1: *Register Access-Use Intervals*. Data results for the BZIP2, COMPRESS, CRAFTY, GCC, and GO programs are shown. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

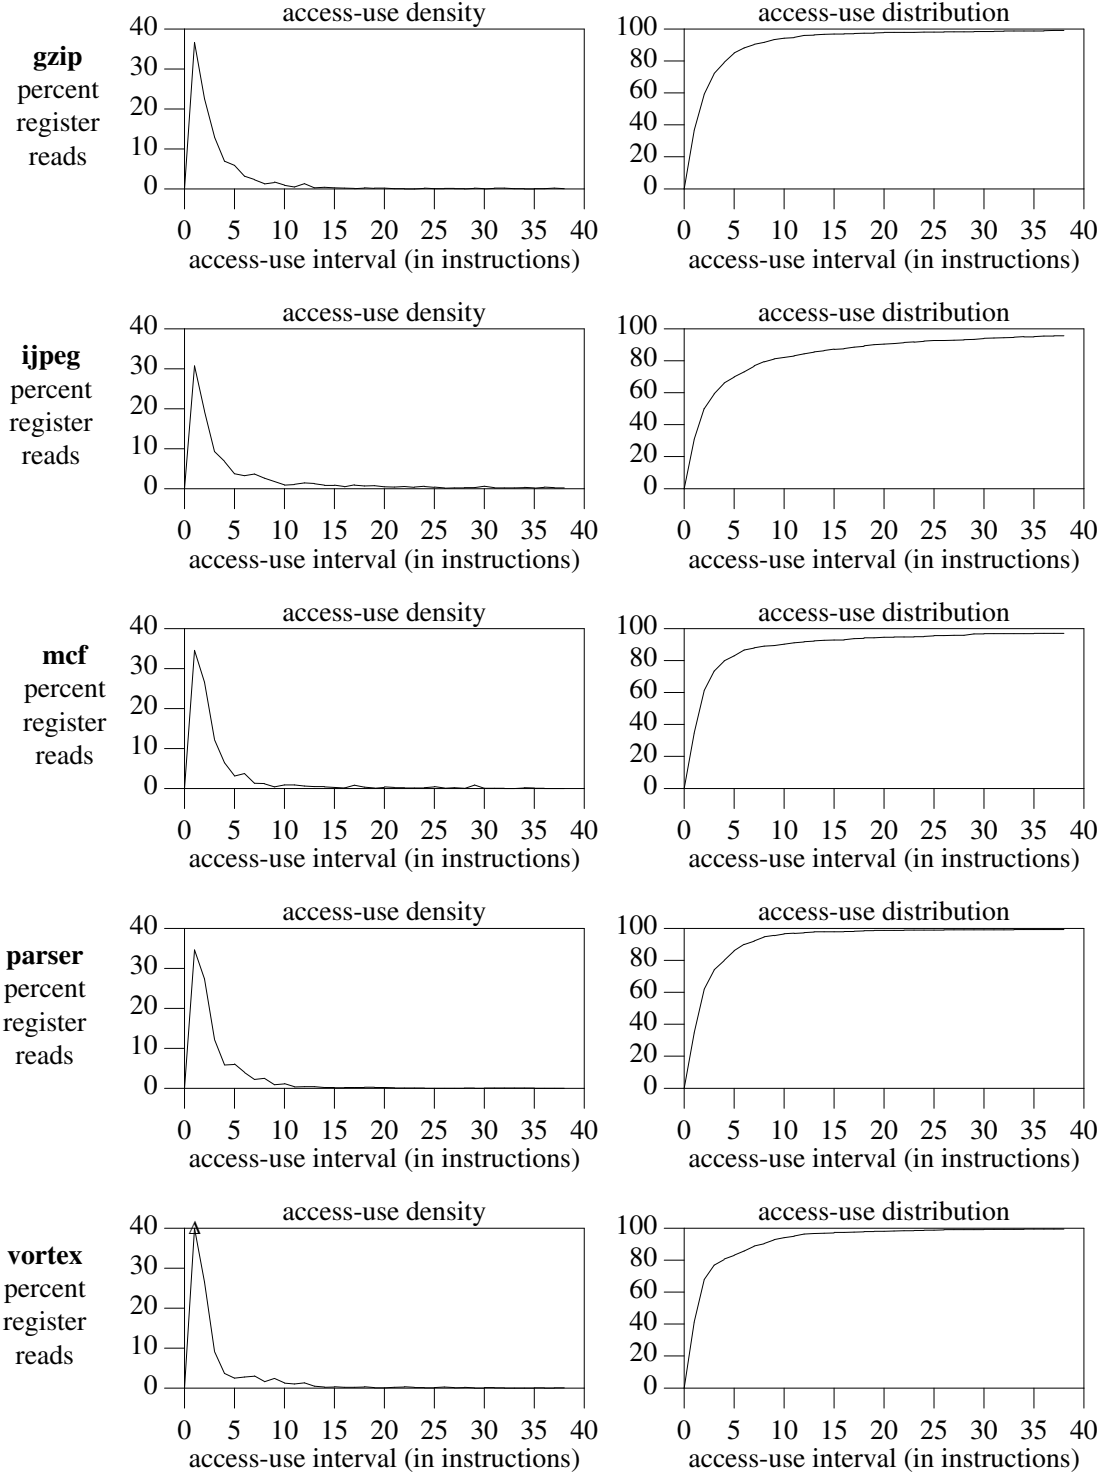


Figure 2: *Register Access-Use Intervals*. Data results for the GZIP, JPEG, MCF, PARSER, and VORTEX programs are shown. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

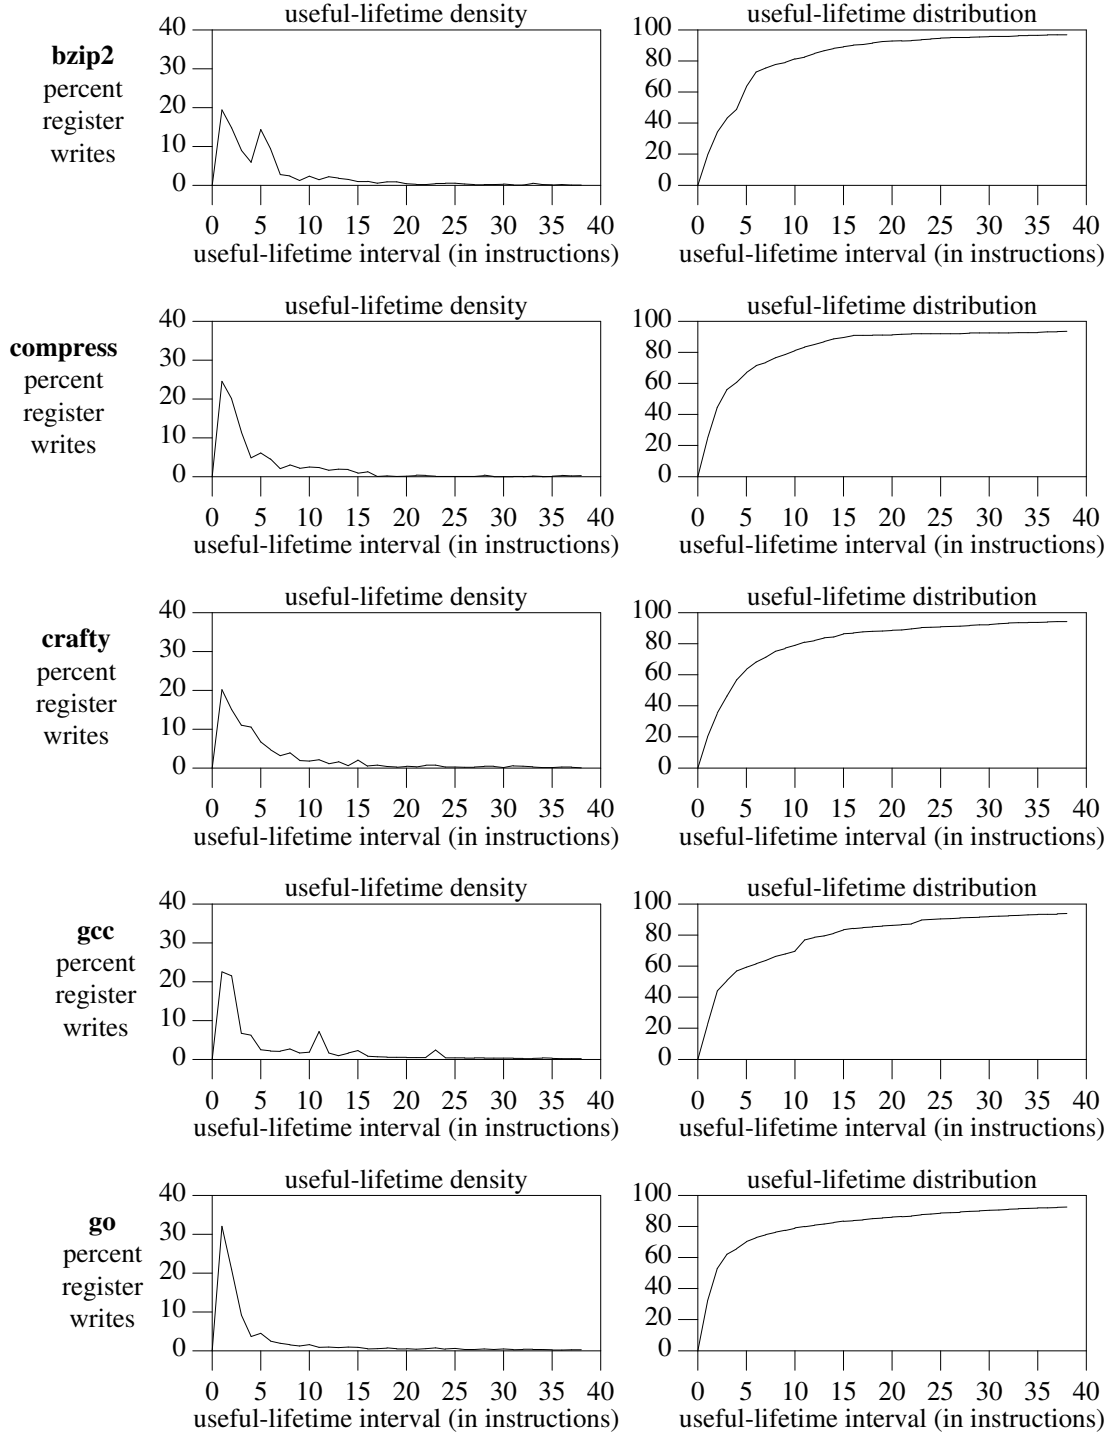


Figure 3: *Register Def-Last-Use Intervals*. Data results for the BZIP2, COMPRESS, CRAFTY, GCC, and GO programs are shown. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

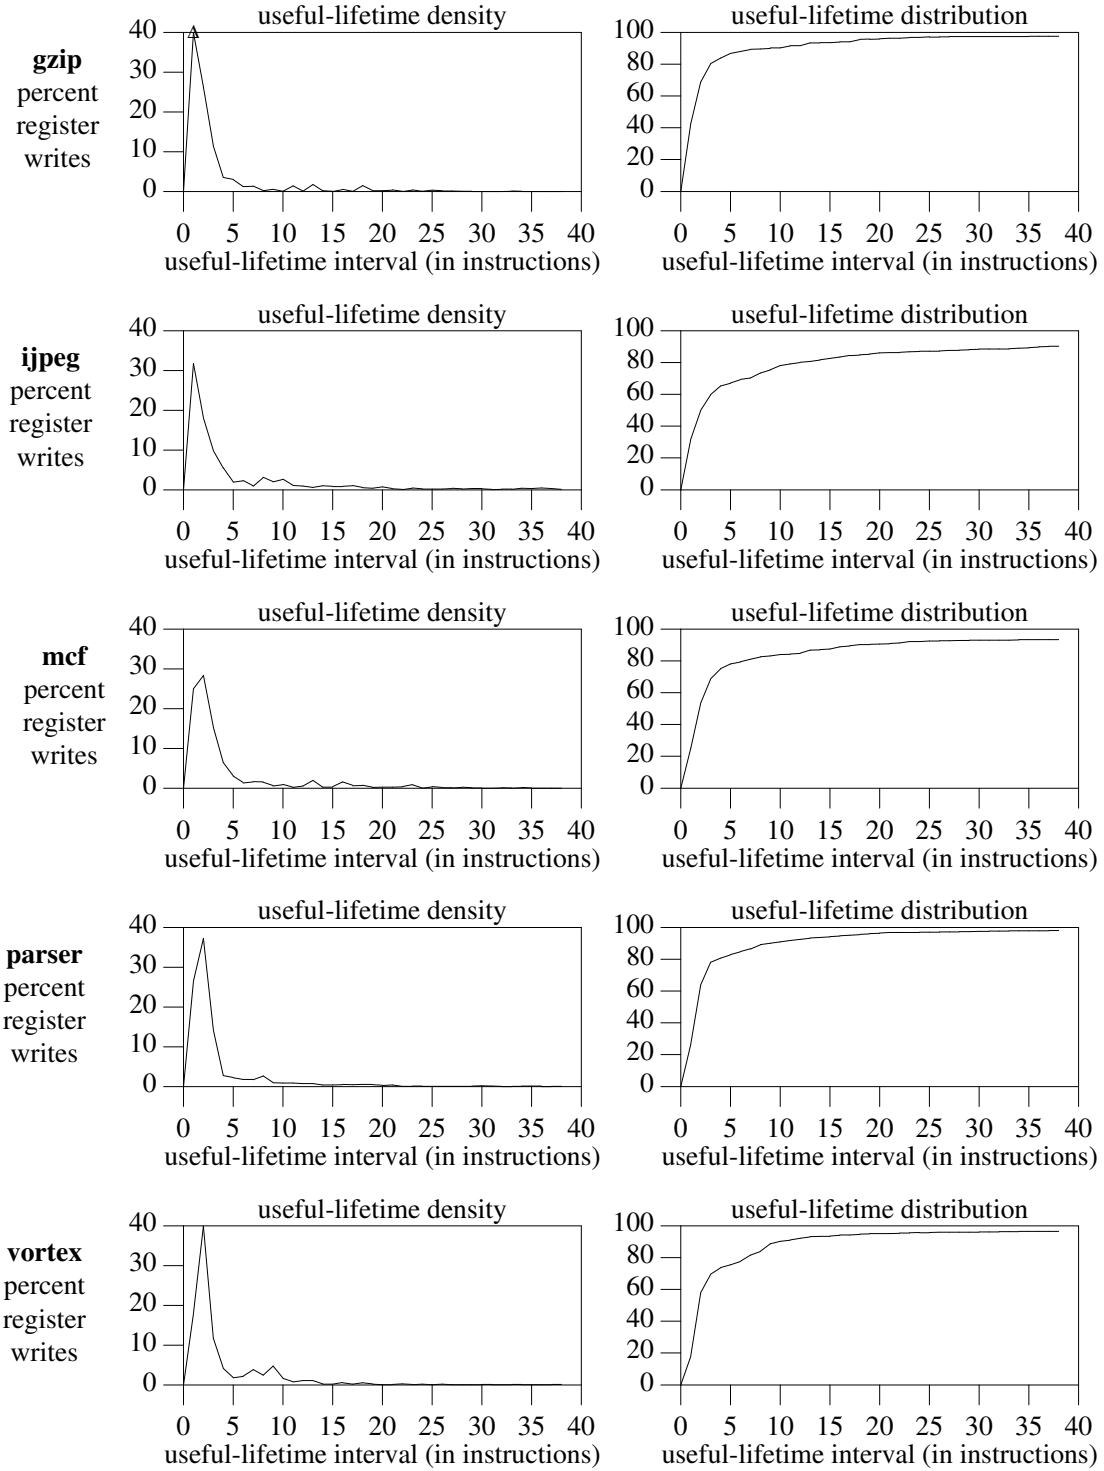


Figure 4: *Register Def-Last-Use Intervals*. Data results for the GZIP, JPEG, MCF, PARSER, and VORTEX programs are shown. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

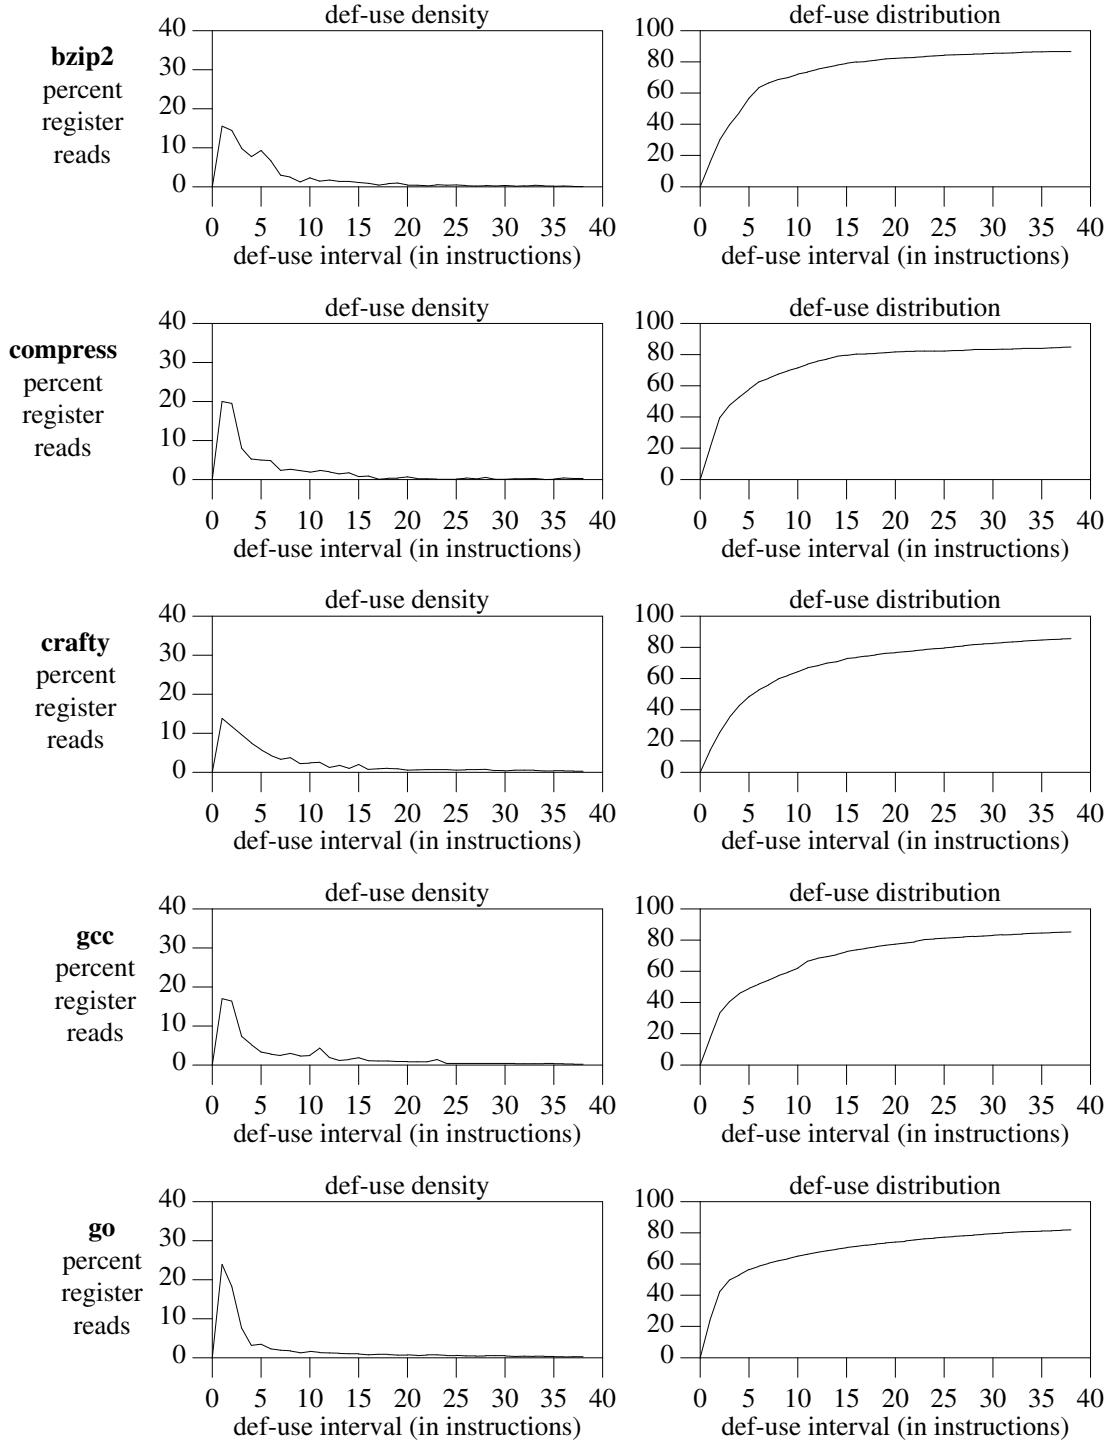


Figure 5: *Register Def-Use Intervals*. Data results for the BZIP2, COMPRESS, CRAFTY, GCC, and GO programs are shown. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

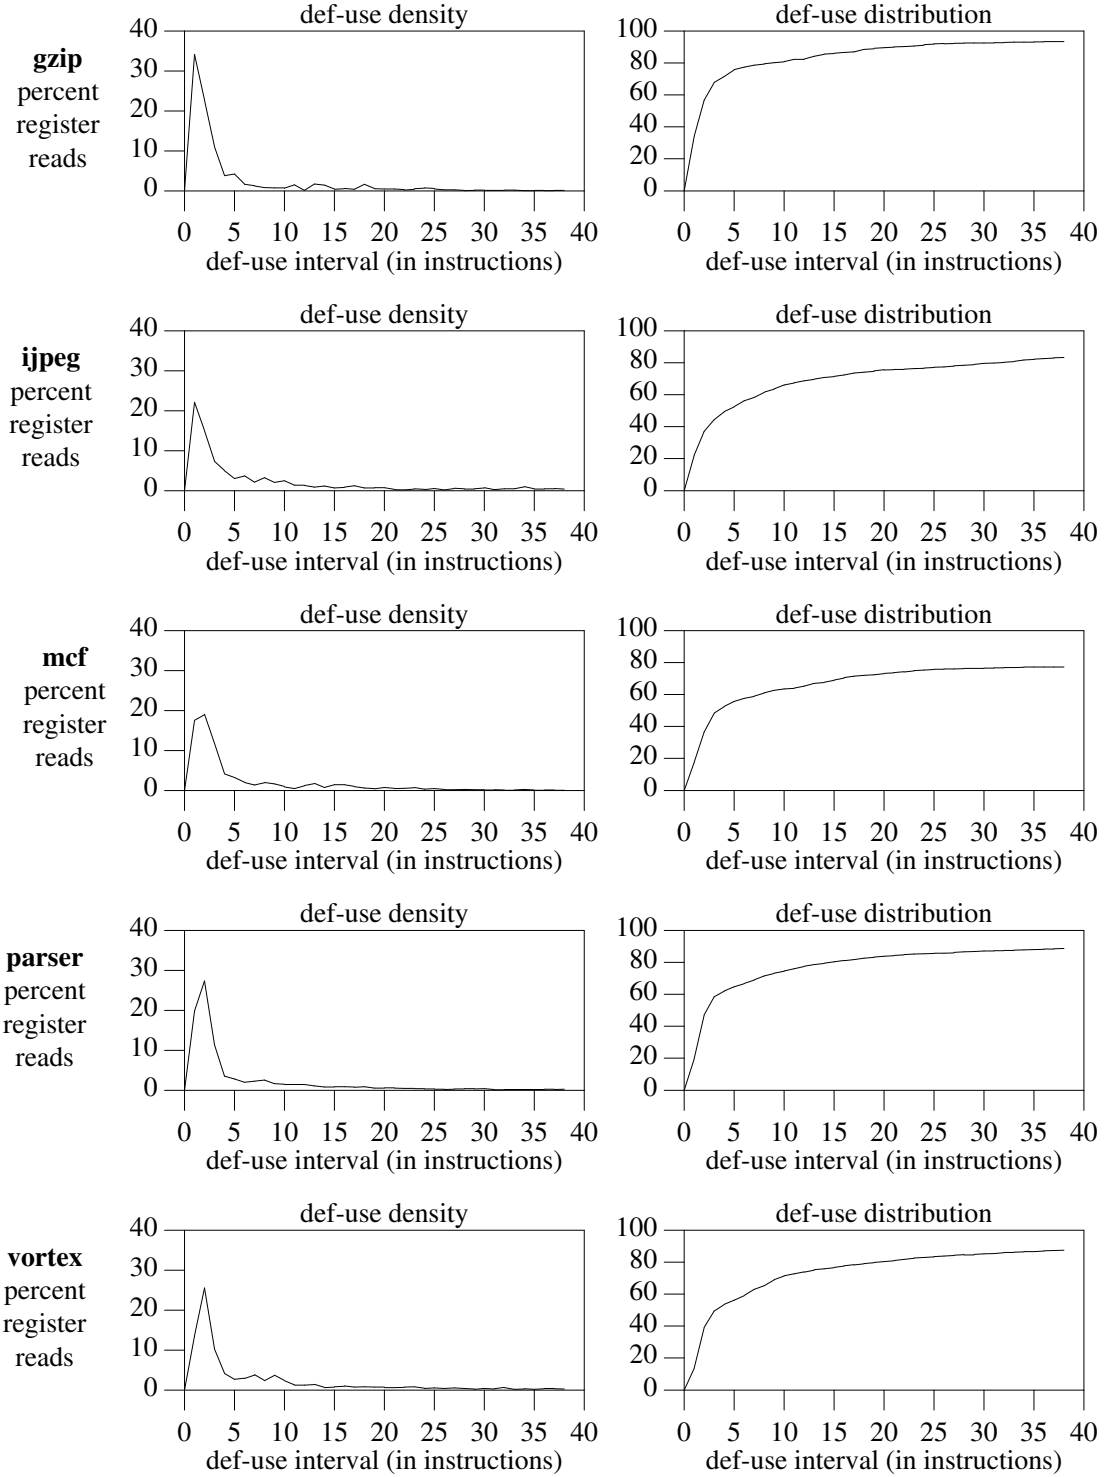


Figure 6: *Register Def-Use Intervals*. Data results for the GZIP, JPEG, MCF, PARSER, and VORTEX programs are shown. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

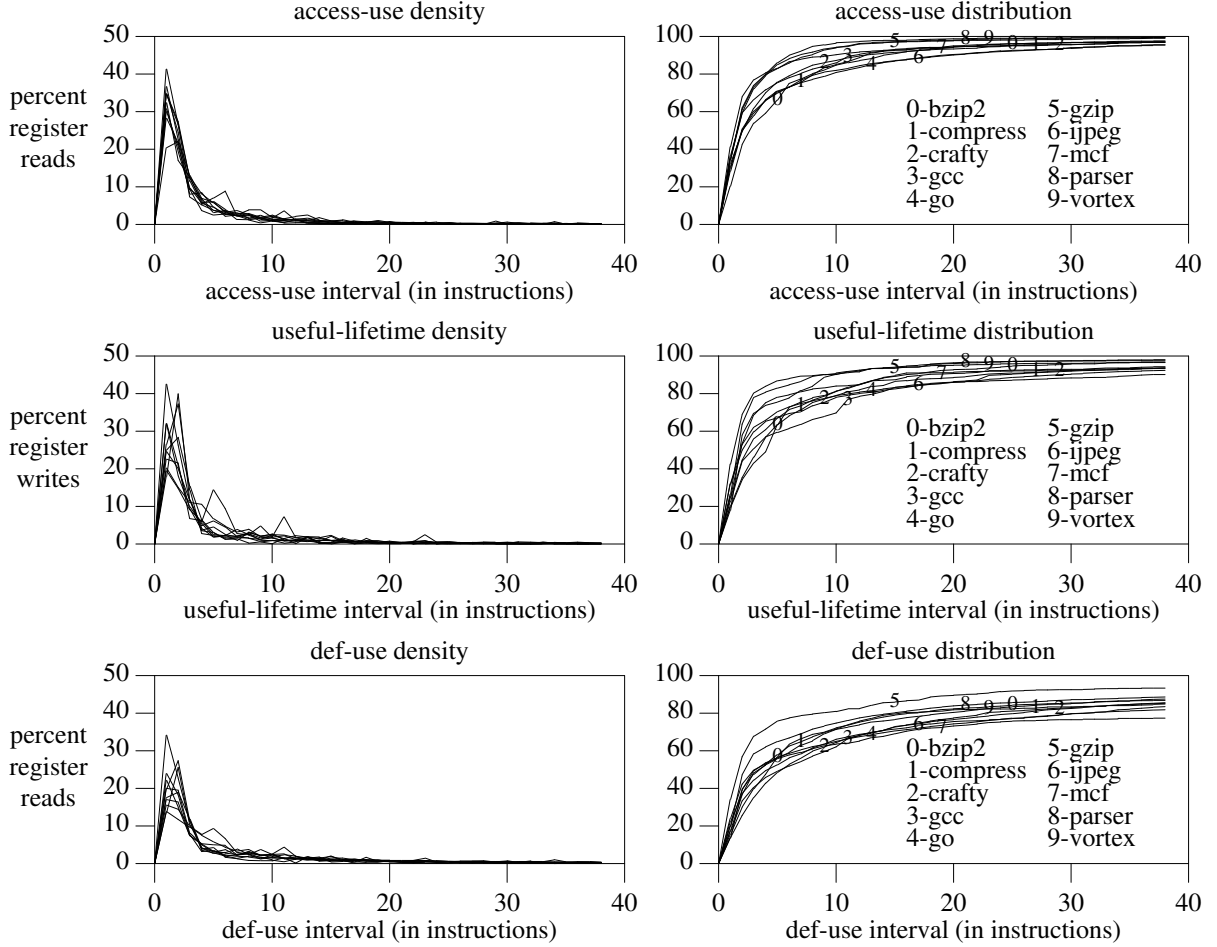


Figure 7: *Register Access Intervals*. Data results for all programs are shown overlaid. The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

can be seen from the graphs of Figure 7, all of the benchmark programs perform similarly and have rather short (20 or less) interval lengths for most of each type of access interval.

For better clarity, in Figure 8 we show the cumulative register interval data over all benchmark programs. That figure shows, in order, all three of the access intervals that we explored: access-use, useful-lifetime, and def-use.

3.2 Register Access Frequency Results

In this section we show the frequencies for reads and writes for each individual register. Each register is identified by its architected address. Separate accounting was maintained for reads and writes so the corresponding data is also separated.

The data for register access frequencies verses register address are shown in Figures 9 and 10. The frequencies for reads and writes are each normalized to all of the register reads and writes respectively. This essentially is a density of register accesses with respect to register address. Results from benchmark programs BZIP2, COMPRESS, CRAFTY, GCC, and GO are shown in Figure 9 while the results from programs GZIP, LJPEGE, MCF, PARSEK, and VORTEX are shown in Figure 10. Figures 11 and 12 show the cumulative




Figure 8: *Cumulative Register Intervals over all benchmarks.* The density is shown on the left and the distribution is shown on the right. All intervals are measured in dynamic numbers of executed instructions.

densities over register addresses for register reads and register writes respectively. On the MIPS ISA, registers $r0$ through $r31$ are integer registers. Registers $r64$ through $r85$ are floating point registers. Other registers (which do not get much use) are various status or control registers (usually for floating pointer operations). It should be noted that register $r0$ is hardwired to the value zero and is a read-only register. For reference, register $r29$ is the stack pointer. Further, register $r28$ is the pointer to the global offset table and is only written once at the start of the program. The data shown is only for the execution of 500 million instructions after skipping the first 100 million, so the single write of register $r28$ would not show up in this data anyway. Registers $r26$ and $r27$ are never written or read by any of the benchmark programs investigated (the ten that we executed). Finally, register $r39$ (visible in some of the graphs) is actually the sum of all register accesses for MIPS registers with address 39 or higher. As a result, this is essentially the sum of all accesses to floating point data or control/status registers.

4 Discussion

Examination of the various register access intervals for each program shows rather similar behavior across most all programs. However, when examining the register access frequency data for each program, shown in Figures 9 and 10, it is evident that most programs make use of the architected registers in noticeably different ways. Although allocation of architected registers is largely done by the compiler (all programs were written in the C programming language), some registers have special standardized meanings as mentioned in the previous section.

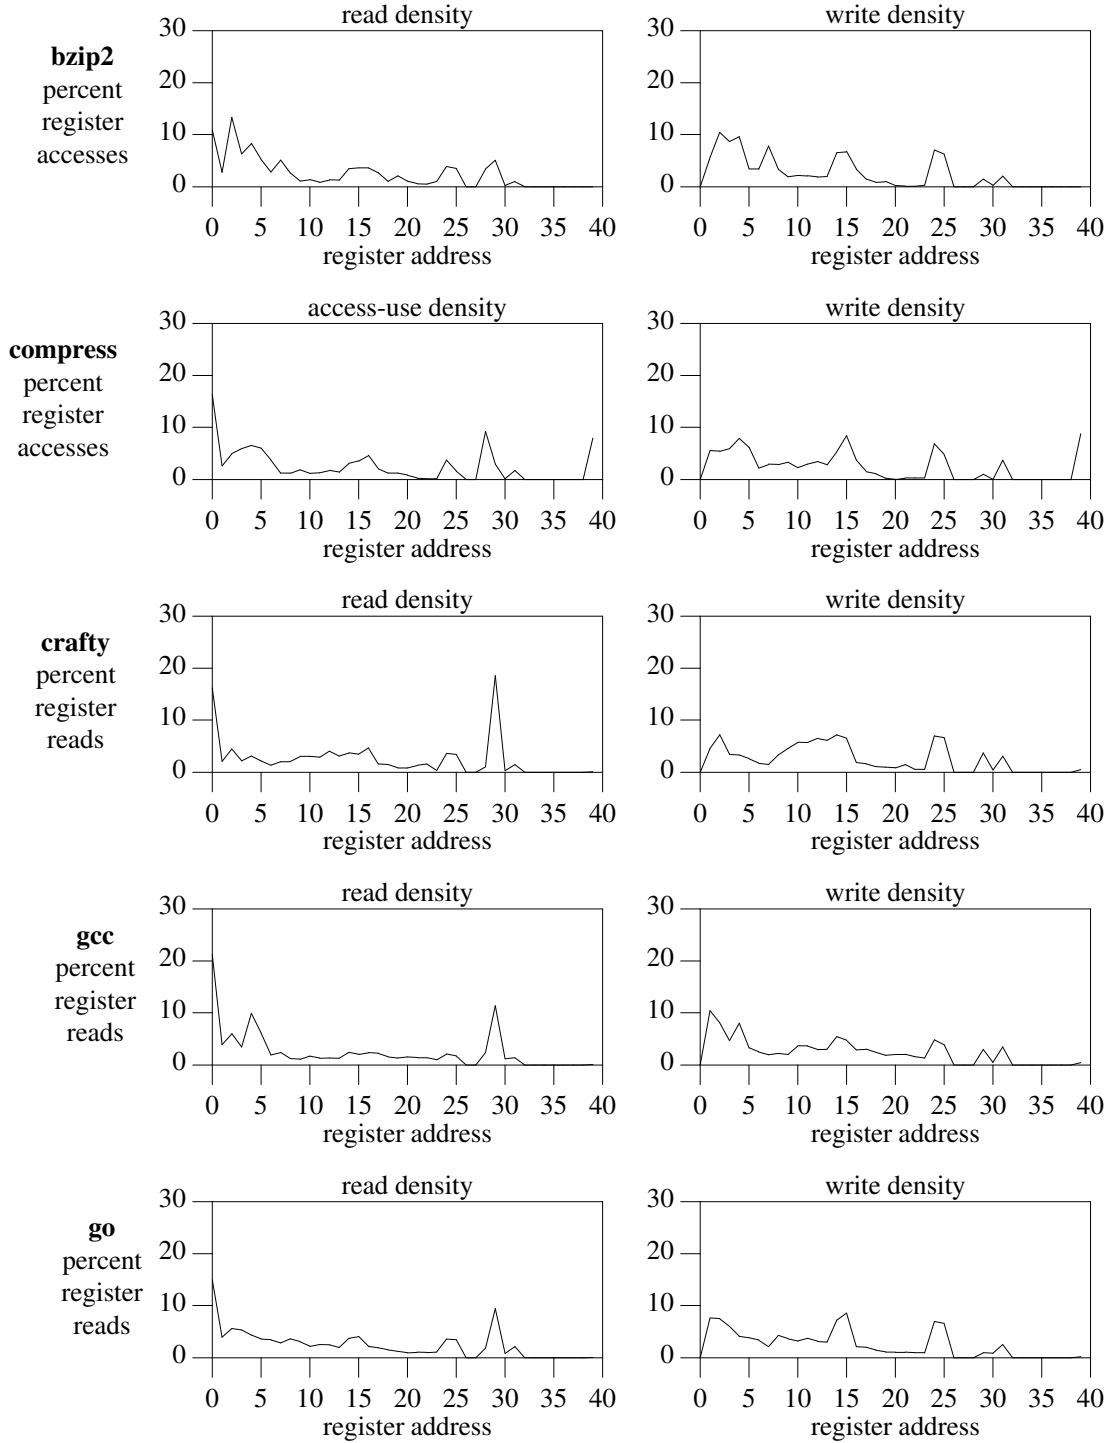


Figure 9: *Register Read and Write Access Frequencies*. Data results for the BZIP2, COMPRESS, CRAFTY, GCC, and GO programs are shown. The read frequencies are shown on the left and the write frequencies are shown on the right. All accesses are shown at specific register addresses.

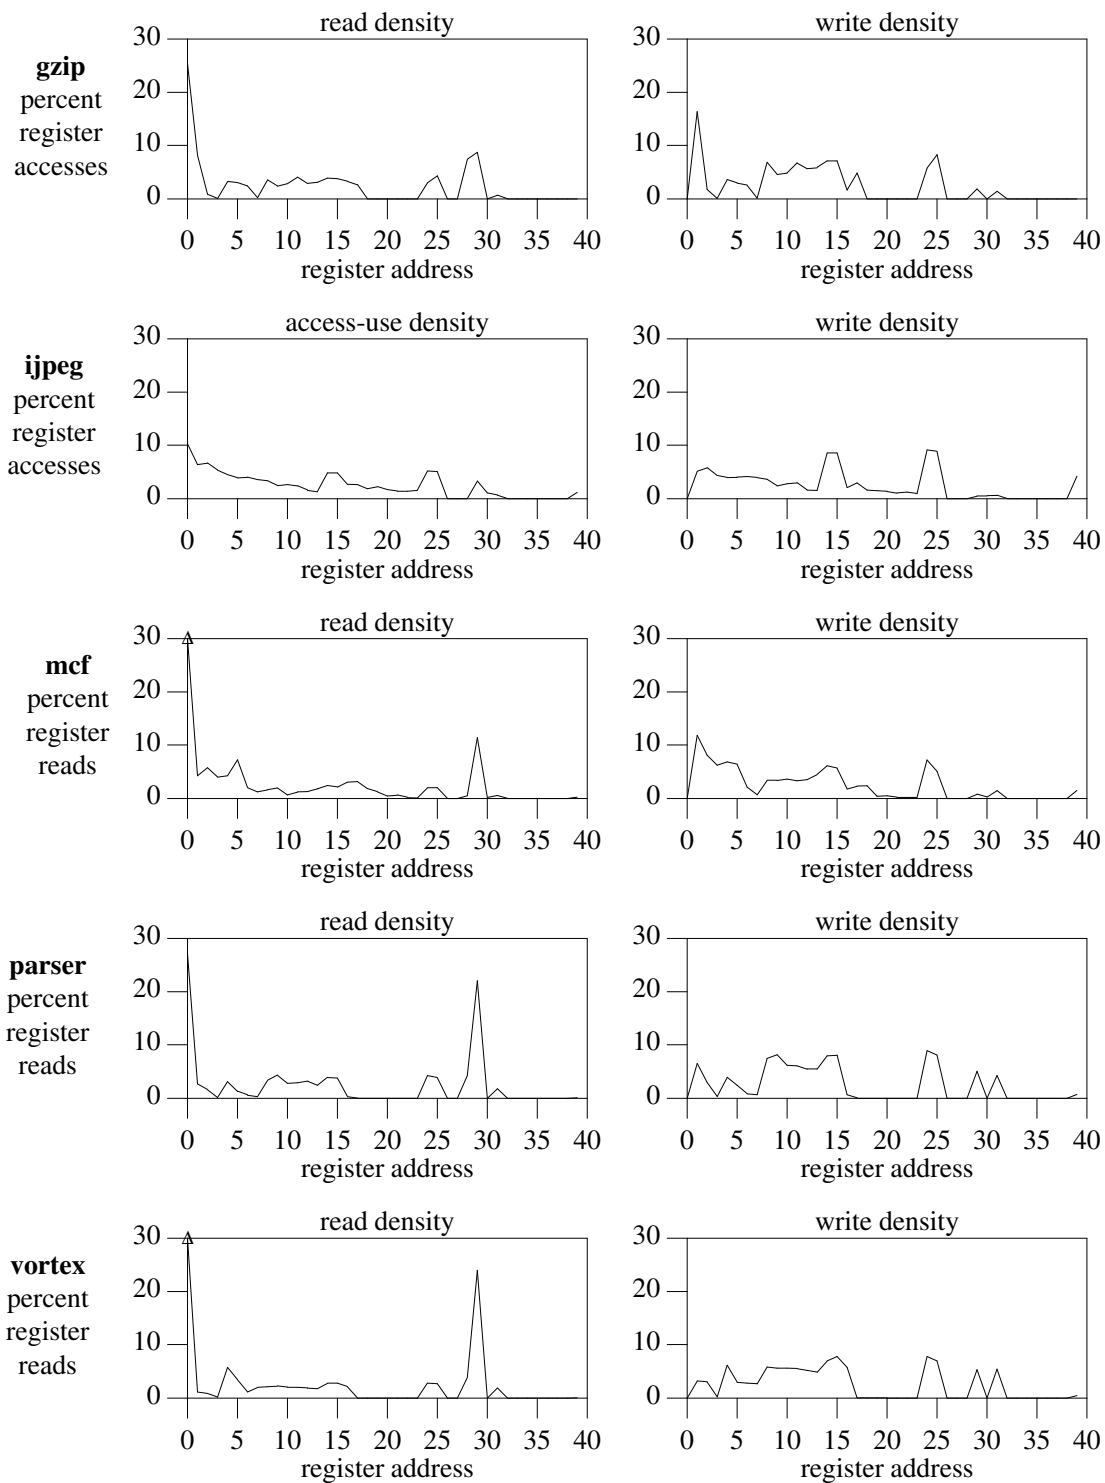


Figure 10: *Register Read and Write Access Frequencies*. Data results for the GZIP, JPEG, MCF, PARSER, and VORTEX programs are shown. The read frequencies are shown on the left and the write frequencies are shown on the right. All accesses are shown at specific register addresses.




Figure 11: *Cumulative Register Read Frequencies*. The cumulative register read frequencies over all benchmark programs are shown. All accesses are shown at specific register addresses.




Figure 12: *Cumulative Register Write Frequencies*. The cumulative register write frequencies over all benchmark programs are shown. All accesses are shown at specific register addresses.

While all programs make use of their memory stacks, typified by reads and writes to register $r29$, some make much more use of it than others. The BZIP2 and IJPEG program make relatively low use of their stack as compared with the others. In contrast, the programs CRAFTY, PARSER, and VORTEX make especially high use of their stacks. Another difference among the programs is the way they make use of register $r0$ (fixed in the MIPS ISA to the value zero). The GZIP, MCF, PARSER, and VORTEX program make especially high use of that while the others make more moderate or smaller use. No programs write register $r0$ since it is fixed as read-only in the ISA. Although all of the benchmark program investigates are "integer" oriented programs, some make substantial use of the floating point register anyway. The COMPRESS, and IJPEG programs make noticeable use of the floating point registers while the other programs make either a small or entirely negligible use of them. Interestingly, the COMPRESS program reads the floating point registers almost as much as it writes them, but this is not the case for the IJPEG program (more on this later). In the case of IJPEG, it writes floating point register much more than it reads them. This generally indicates that conditional control flow is abandoning writes that occur before conditional branches. Although this contributes towards its register useful-lifetimes being zero for these cases (reference Figure 4), the effect does not appear to be significant as compared with the other benchmarks.

It is interesting to note generally that the register read frequencies are quite similar for the GCC, GO, and MCF programs. This is also the case for the CRAFTY, PARSER, and the VORTEX programs. The PARSER and VORTEX programs are especially similar even though the purpose of each program is different (PARSER performs word processing and word indexing while VORTEX is a database kernel). For write frequencies, they appear to be largely more similar across the programs than was the case for read frequencies. The notable exceptions are CRAFTY, GZIP, PARSER, and VORTEX all having more register writes in the 8 to 13 range of register address than the others. GZIP also has a singularly large percentage of write to register $r1$ than any of the others.

Another important feature of these graphs is that write frequencies (for almost all programs) do not match the corresponding read frequencies of the same program. This behavior is especially evident around the program memory stacks, which use architected register $r29$. However, most registers are written with different frequencies than they are read. Although it is generally expected that one write to a particular register might be read more than once, the case where more than one write to a register with only one or less succeeding read also occurs. These cases are largely due to the register write being abandoned due to intervening conditional control flow changes. The different in register read and writes frequencies is part of the reason why register useful-lifetimes are normally so different from register def-use intervals.

If, for example, all registers were written and then followed by one read before they were written again, then the register useful-lifetime intervals and the register def-use intervals would be identical. But this is generally not the case. Rather, some registers are written more often than read (giving rise to smaller average useful-lifetimes) while the reverse can also occur (giving rise to average smaller def-use intervals). Further, even if the cumulative read and write frequencies were the same in any given program, we still do not have information on particular read and write instances (only that the average frequencies are the same). These and other effects generally make all three types of register access intervals investigated to be different.

5 Summary

We have presented supplementary data on register access intervals that was not presented in the previous paper. Three types of intervals were explored. There were: access-use, useful-lifetime, and def-use intervals. Additionally, we presented data for register read and write access frequencies.

References

- [1] Khalafi A., Morano D., Kaeli D., Uht A. Realizing High IPC Using Time-tagged Resource-Flow Computing. In *submitted for the ISPASS Conference 2003*, 2003.